
Computer networks and the Internet — TTM4200

A. Murad, D. Palma, S. Fuglesang, C. Lewin

Contents

Introduction	3
Setting up Virtual Machine	4
1 Milestone 1 – Getting Started with Docker	6
1.1 Creating Docker Images	7
1.2 Running Containers	7
1.3 Docker Network	8
1.4 Publishing Container's Ports	9
1.5 Docker Volumes	9
2 Milestone 2 – Docker Compose	9
2.1 Creating Docker Compose Files	10
2.2 Running Docker Compose Files	12
3 Milestone 3 – Traffic Capture & Analysis Tools	12
3.1 Terminal multiplexer (tmux)	12
3.2 Traffic Capture	14
3.3 Traffic Analysis	15
4 Optional Exercise	16

Introduction

The goal of this lab is to experiment with the main tools we are going to use in this course, including Docker, Docker Compose, Tcpdump, Wireshark and Tmux. The lab has several **milestones**. Make sure you reach each one before advancing to the next.

For delivery, submit a PDF report where you answer **only** those steps that are marked with **REPORT(%)**. Additionally, you should submit the codes or capture files, if they were **explicitly** asked for. The percent point gives you an indication of the score of that question. Lab1 counts for **2 points** of your final score in this course.

The answers should be clear and explicit. First, write the section number and title, then question number and the question itself, and then your answer. The screenshots should be annotated. You should provide a single screenshot for each question, i.e. do **not** refer to a single screenshot for multiple questions. An example of an excellent answer:

3.3 Traffic Analysis

Q3. How long did it take from when the **ping** request was sent **until** the reply was received?
A3. According to the reply packet, the response **time** was 0.031ms, as shown **in** the following screenshot.

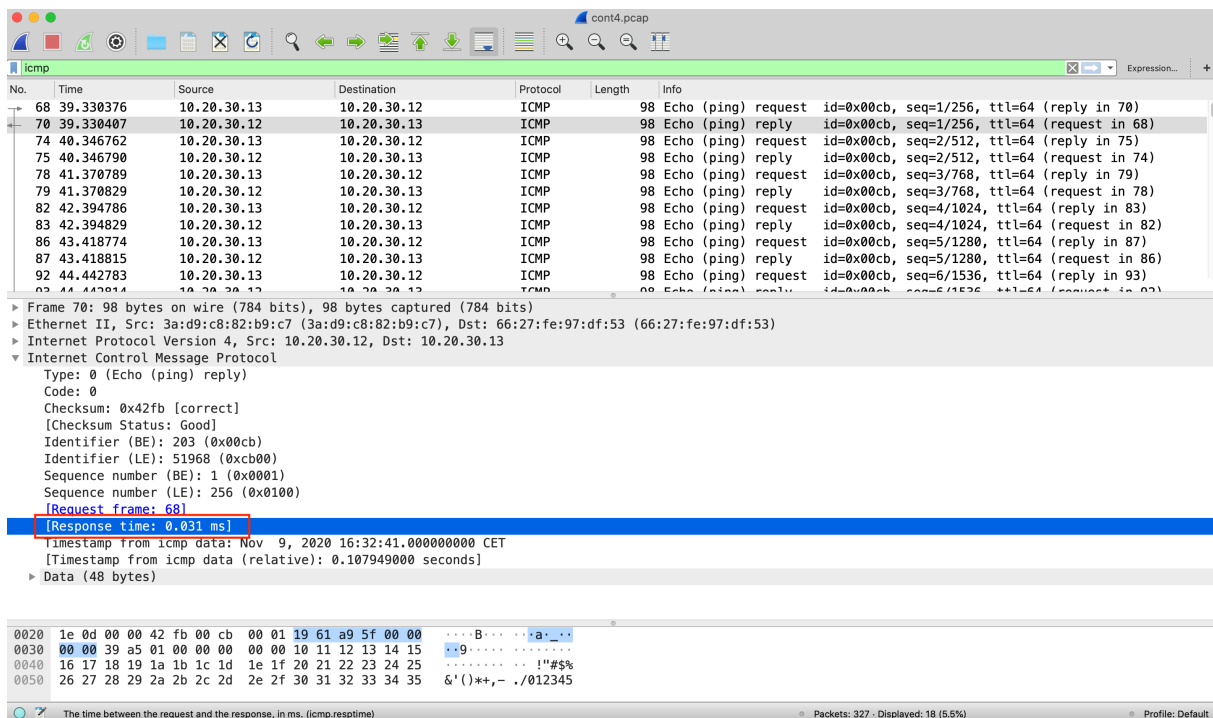


Figure 1: Annotated Screenshot

We will score your answers according to a rubric, for example:

Grading NO	Question	Model answer	Excellent Answer (100%)	Good answer (80%)	Needs improvement (40%)	Band zero (0%)	Grade (%)
3.3 Traffic Analysis							
Q3	Q3. How long did it take from when the ping request was sent until the reply was received?	A3. According to the reply packet, the response time was *** ms.	Close to the model answer and an annotated screenshot	Close to the model answer. and screenshot.	Includes only a relevant screenshot without the value or includes just a related value without a screenshot.	Gave a random answer without a screenshot, or random screenshot.	5%

Figure 2: Rubric for Q3

Setting up Virtual Machine

- Start by downloading and importing the VM ([Link](#)) and configure it according to your computer's specifications (as you did in TTM4175. Additionally, to access files of your computer from within the VM, you should share folders between them. Figure 3 summarizes the process of setting up the VM (adapt the settings to your own case).
- Start the VM and login using the password “ttm4200”. Do not upgrade if prompted to upgrade to newer Ubuntu version. Change the password if you haven't done so (`passwd`).
- By default, the shared folder will be mounted to “/media/sf_shared” (if Auto-mount is selected). Add the user “ttm4200” to “vboxsf” group, then reboot your VM:

```
sudo usermod -a -G vboxsf ttm4200
sudo reboot
```

You can mount the shared folder in any directory of your choice ([random guide](#)).

- Before starting lab1, you are encouraged to remove any containers and networks that you may have built before in your VM. This is to prevent unnecessary conflicts later in this lab:

```
docker rm $(docker ps -aq)
docker network prune
```

- From Blackboard, download “lab1.zip” to your VM and place it in your home directory (“/home-/ttm4200”). Unzip this folder and navigate to it:

```
cd /home/ttm4200
unzip lab1.zip && cd lab1
```

All the provided commands in this lab assume that your present working directly (`pwd`) is “/home/ttm4200/lab1”, and they won't work if you are inside another directory.

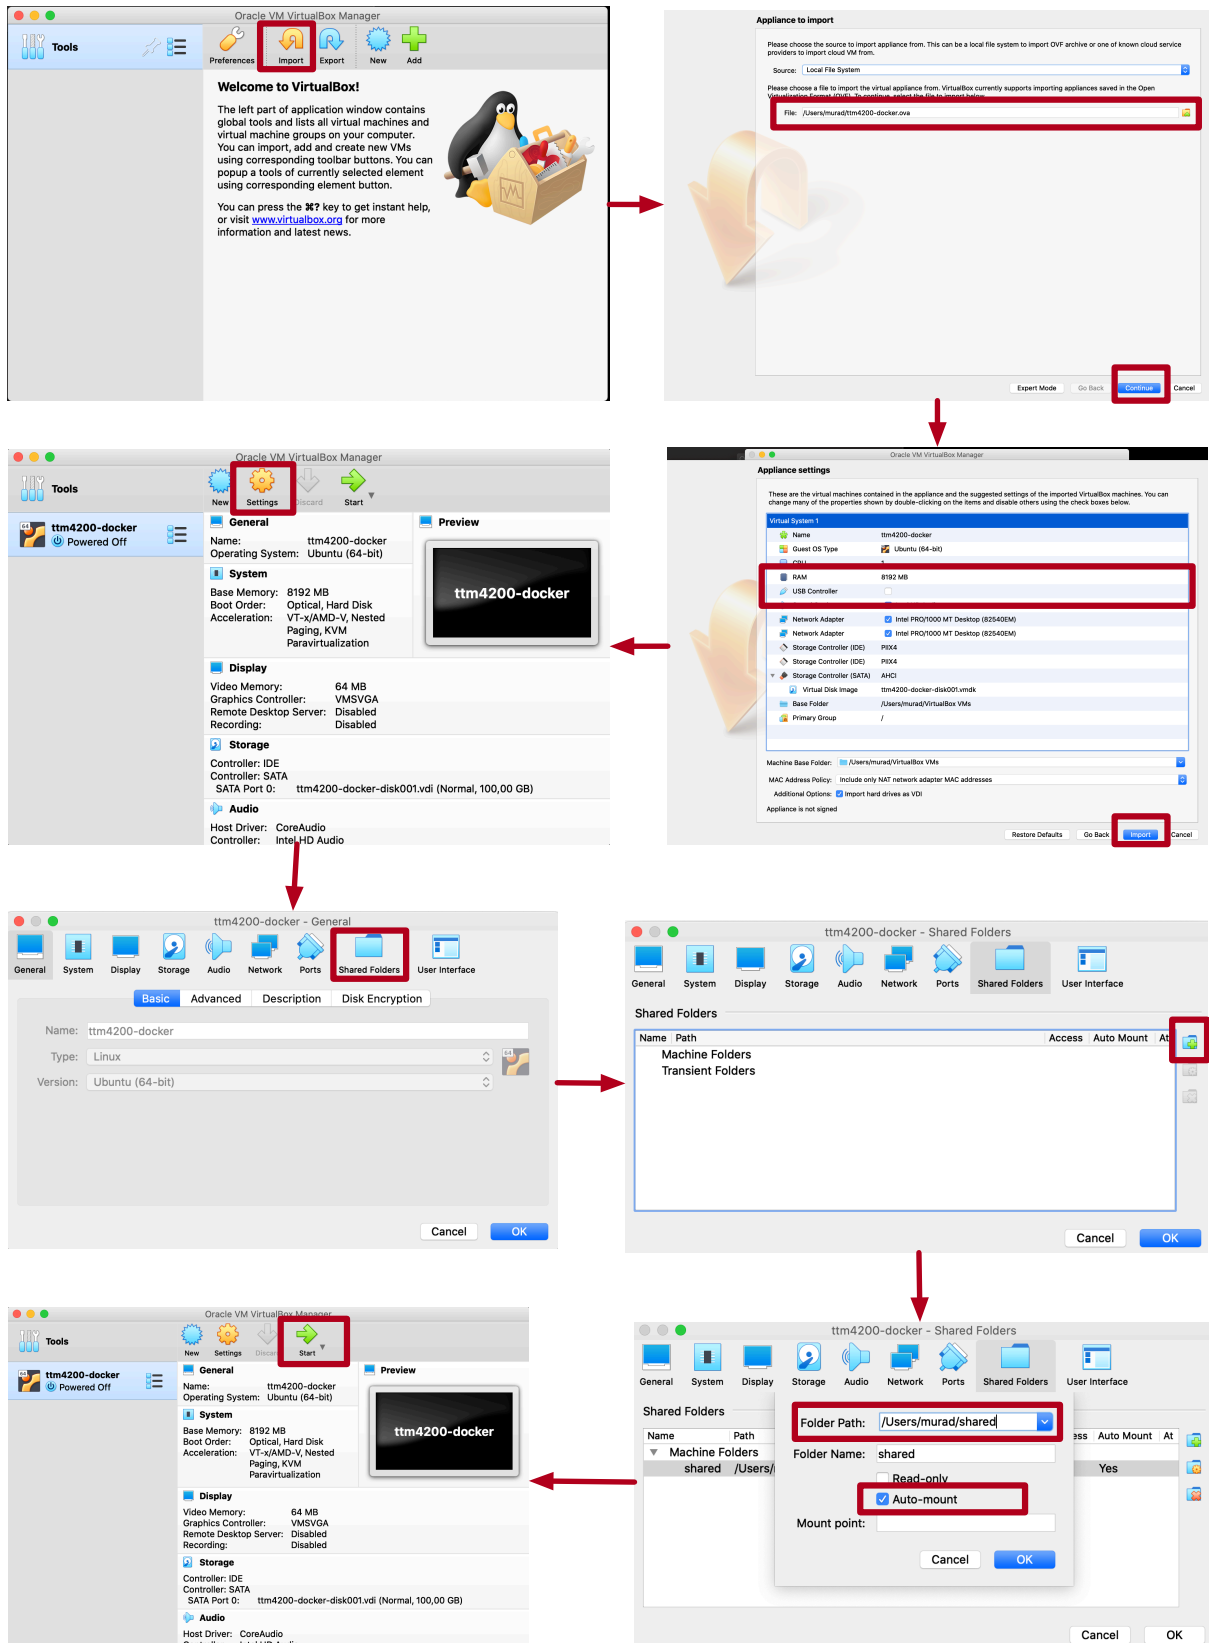


Figure 3: VM Set up

1 Milestone 1 – Getting Started with Docker

In this lab, you should get acquainted with [Docker](https://docs.docker.com/) (<https://docs.docker.com/>), since we will use it extensively in all future labs. Read this [overview](https://docs.docker.com/engine/docker-overview/) (<https://docs.docker.com/engine/docker-overview/>) as well as the support document to get an idea of what Docker is.

By the end of this milestone, you will build a system as shown in Figure 4 and Figure 5. Refer back to these figures as often as needed to understand how the lab is constructed.

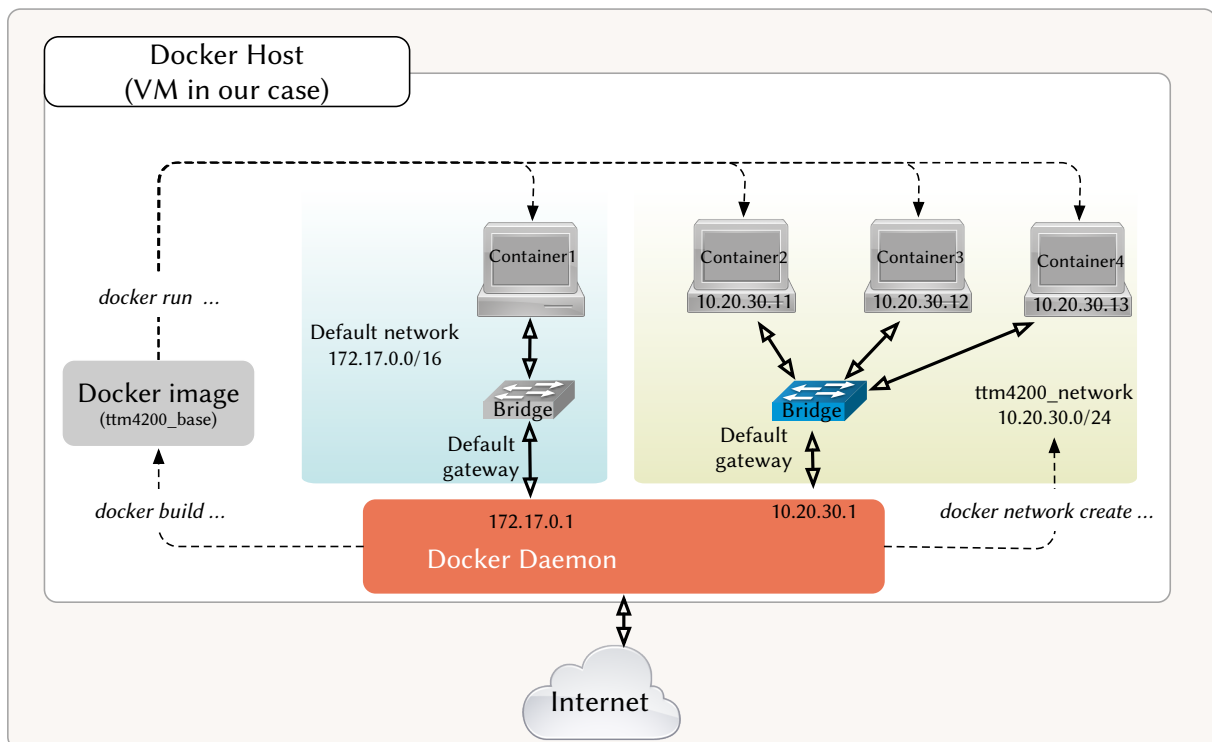


Figure 4: Docker Image, Containers, and Networks

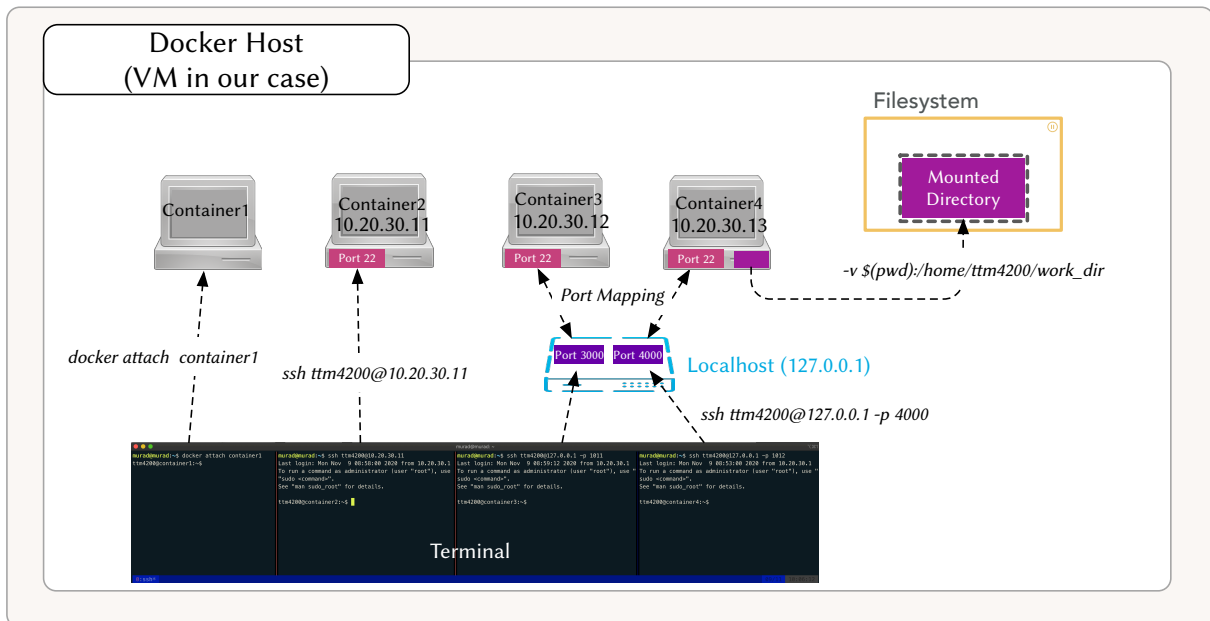


Figure 5: Docker Attach, SSH, Port Mapping, and Volumes

1.1 Creating Docker Images

To build a docker image, we need to write a [Dockerfile](https://docs.docker.com/engine/reference/builder/) (<https://docs.docker.com/engine/reference/builder/>) that contains all the commands for assembling an image. In this course, all Dockerfiles will be provided to you. We will always build our images on top of [Ubuntu](https://hub.docker.com/_/ubuntu/) (https://hub.docker.com/_/ubuntu/).

- Build an image from the Dockerfile in the directory named “ttm4200_base”:

```
docker build --tag ttm4200_base ./ttm4200_base/
```

- Check how many images are built locally:

```
docker images
```

1.2 Running Containers

- Run a container and name it “container1” from the image you have just built. Run the container interactively, and in the background (detached mode):

```
docker run --interactive --tty --detach --name=container1 --hostname=container1 \
    ttm4200_base
```

- Attach a terminal to the container. This will allow you to control the container interactively in your terminal.

```
docker attach container1
```

The password for the root user inside the container is set to “ttm4200”

- Check the container IP address, MAC address, routing table, and ARP table.

```
ip address
ip link
ip route
ip neighbor
```

By default, Docker uses the 172.17.0.0/16 subnet for container networking, and assigns the first address (172.17.0.1) to the host (the machine which is running Docker daemon, the VM in your case).

- Ping the Docker host, then again check the ARP table.
- Detach from the container (press `Ctrl+p` followed `Ctrl+q`). This will turn the container from interactive mode into daemon mode.

If you press `Ctrl+c` or `exit`, you will kill the container. This is one of the common pitfalls (especially in later labs)

1.3 Docker Network

- Create a new docker network with a specific name and subnet:

```
docker network create --subnet=10.20.30.0/24 ttm4200_net
```

- Run a second container (name it “container2”) and connect it to the new network, and assign it a static IP address:

```
docker run -itd --name=container2 --hostname=container2 --net=ttm4200_net \
--ip=10.20.30.11 ttm4200_base
```

- Connect to the container using SSH. The ssh server is already installed and configured in the Dockerfile. It is configured to accept only a public key authentication to the user “ttm4200”, and the private key is already in the VM (“/home/ttm4200/.ssh”). You can connect directly using the container ip address:

```
ssh ttm4200@10.20.30.11
```


1.4 Publishing Container's Ports

By default, a container does not make its port available to the outside world (or the Docker host). To do so, you have to map a container port to a port on the Docker Host.

- Run a third container (“container3”) and bind its SSH port to an unused port in your Docker Host (preferably a number above 1024, to avoid conflicts), and make it listen on localhost only (127.0.0.1).

```
docker run -itd --name=container3 --hostname=container3 --net=ttm4200_net \
  --ip=10.20.30.12 -p 127.0.0.1:3000:22 ttm4200_base
```

if you don't specify 127.0.0.1 in the port mapping, it will publish the port on all interfaces of the Docker Host (0.0.0.0). This is less ideal from a security perspective.

- Connect to the container using SSH, but using the published port on the localhost of Docker Host:

```
ssh ttm4200@127.0.0.1 -p 3000
```

1.5 Docker Volumes

To preserve data generated by a container, you can either copy them to your host machine or use *Volumes* for persisting data. A Docker volume is basically a mechanism to mount a file system (directory) between a Docker host and a container. Thus, any data generated will be stored automatically in your host machine.

- Run a fourth container (“container4”) and mount the “ttm4200_base” directory to a directory in the container:

```
docker run -itd --name=container4 --hostname=container4 --net=ttm4200_net \
  --ip=10.20.30.13 -p 127.0.0.1:4000:22 \
  -v $(pwd)/ttm4200_base:/home/ttm4200/work_dir ttm4200_base
```

- Stop and remove all the containers as well as the network.

```
docker stop $(docker ps -aq)
docker rm $(docker ps -aq)
docker network prune
```

2 Milestone 2 – Docker Compose

In this course, we are going to use Docker extensively to build several different services using containers. For handling multiple containers at once, we need to use Docker Compose. In Docker Compose, we will create a YAML file to configure the application's services, and with a single command, we can

create all the images and run the containers at once. Additionally, we will be able to stop and remove all containers and the network in a single command. For further reference, see the documentation from [docker-compose.yml](https://docs.docker.com/compose/compose-file/compose-file-v2/) (<https://docs.docker.com/compose/compose-file/compose-file-v2/>), and the support document.

2.1 Creating Docker Compose Files

- Complete the “docker-compose.yml” file that defines exactly all we did in Milestone 1 (see Figure 4 and Figure 5), specifically:
 - Build an image, called “ttm4200_base”, from the Dockerfile in the “ttm4200_base” directory
 - Create a network, called “ttm4200_net”, with the subnet “10.20.30.0/24”.
 - Run “container1” interactively, with a hostname of “container1”, and connected it to the default network.
 - Run “container2” interactively, with a hostname of “container2”, and connected to “ttm4200_net”.
 - Run “container3” interactively, with a hostname of “container3”, connected to “ttm4200_net”, and its SSH port mapped to a port in Docker Host.
 - Run “container4” interactively, with a hostname of “container4”, connected to “ttm4200_net”, its SSH port mapped to a port in Docker Host, and the “ttm4200_base” directory is mounted to “/home/ttm4200/work_dir” inside the container.
 - Additionally, assign IPv6 addresses to containers connected to “ttm4200_net”, within the subnet of “fd00::/64”.
 - To make things easier, the network, building the image, “container1” and “container4” are already configured. You must complete “container2” and “container3”!

```
version: '2.3'
#specify a custom network
networks:
  #the name of the created network
  ttm4200_net:
    enable_ipv6: true
    ipam:
      config:
        - subnet: 10.20.30.0/24
        - subnet: "fd00::/64"

services:

  #Building the base image
  image:
```

```
#Build an image from the Dockerfile in the "ttm4200_base" directory
build: ./ttm4200_base
#Name the image "ttm4200_base", and use it
image: ttm4200_base

container1:
  #use the image "ttm4200_base"
  image: ttm4200_base
  #It will depends the service "image", that built the image
  depends_on:
    - "image"
  #Name of the container "container1"
  container_name: container1
  #Name the host machine "container1"
  hostname: container1
  #Run in "--interactive" mode
  stdin_open: true
  #Run in "--tty", pseudo terminal
  tty: true
  # Add networking capabilities, instead of privileged
  cap_add:
    - NET_ADMIN

container2:
#####
# TODO: Your code here!

#####

container3:
#####
# TODO: Your code here!

#####

container4:
#####
  image: ttm4200_base
  depends_on:
    - "image"
  container_name: container4
  hostname: container4
  stdin_open: true
  tty: true
  cap_add:
    - NET_ADMIN
  #map ssh port to a port on Docker Host (port 400)
  ports:
    - "127.0.0.1:4000:22"
  # mount the "ttm4200_base" directory on Docker Host to the container
  volumes:
```

```
- ./ttm4200_base/:/home/ttm4200/work_dir
#Connect to the network (ttm4200_net and assign IPv4 and IPv6 addresses)
networks:
  ttm4200_net:
    ipv4_address: 10.20.30.13
    ipv6_address: "fd00::13"

#####
```

Docker Compose might use a different default network subnet other than 17.17.0.0/16, depending on if this subnet is already in use

Q1. REPORT(20%): Submit your docker-compose file along with the report. We will check that it is working, and that it corresponds to what is shown in Figure 4 and 5.

2.2 Running Docker Compose Files

- Create and start the containers using the docker-compose file:

```
docker-compose up -d --build
```

The “-d” option is to run containers in the background (detached mode), and “--build” is to build the images before starting containers.

- Connect to the containers, as shown in figure 5.

```
docker attach container1
ssh ttm4200@10.20.30.11
ssh ttm4200@127.0.0.1 -p 3000
ssh ttm4200@127.0.0.1 -p 4000
```

3 Milestone 3 – Traffic Capture & Analysis Tools

3.1 Terminal multiplexer (tmux)

We will use **tmux** (<http://man.openbsd.org/OpenBSD-current/man1/tmux.1>) to access multiple terminal sessions simultaneously in a single window. It lets you switch easily between several programs in one terminal. Additionally, you can detach and re-attach from sessions, thus allowing you to run sessions in the background. This is really helpful if you are working on a remote server. [Here \(https://shortcutworld.com/tmux/linux/tmux_Shortcuts\)](https://shortcutworld.com/tmux/linux/tmux_Shortcuts) are the most common **tmux** shortcuts.

- In “container4”, start a new **tmux** session (preferably a named session to distinguish between multiple sessions):

```
tmux new -s [session_name]
```

- Create four panes as shown in Figure 6.

To execute a command in `tmux`, a prefix must be used before the command (`<prefix> + command`). By default, the prefix is `Ctrl+b`. For example, to create a vertical split use `<prefix> + %`. That is, you have to press the keys `Ctrl+b` followed by the command `%` (without the control key). To create a horizontal split use `<prefix> + "`. To move between panes use `<prefix> + Arrow Keys`.

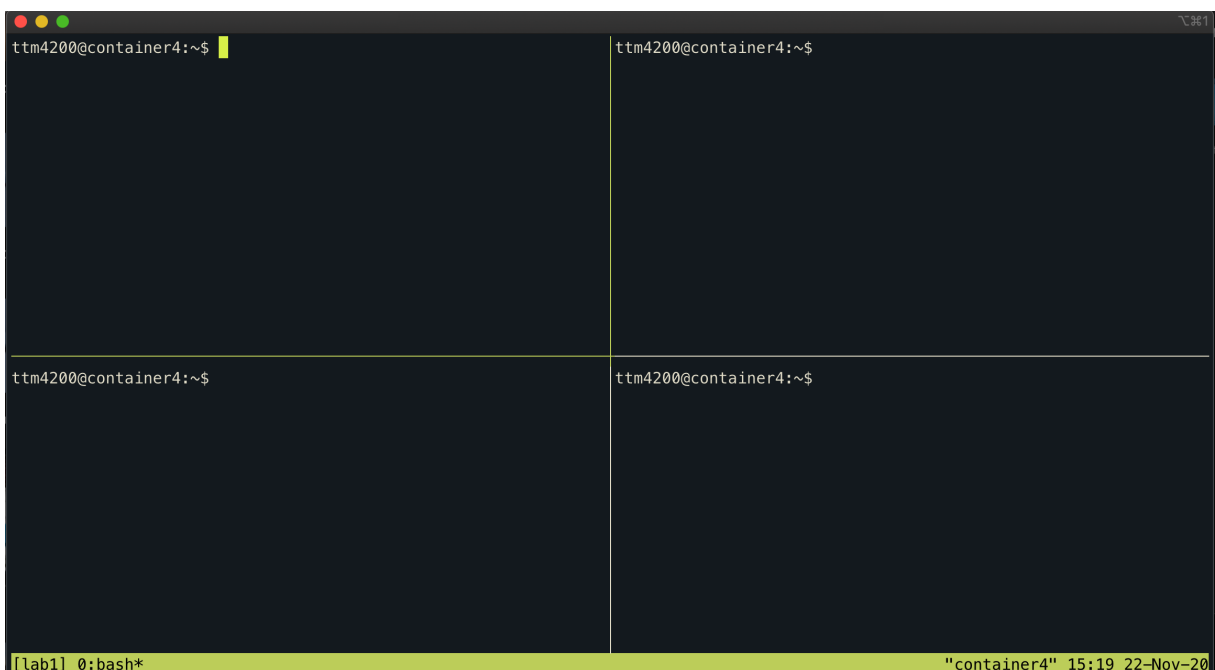


Figure 6: Tmux Panes

- Detach from the current session by using `<prefix> + d`.
- **Optional:** If you feel that `tmux` key combinations are awkward, you can customize it according to your personal preferences. For example, to change the prefix from `Ctrl+b` to `Ctrl+a`, simply add the following lines to `tmux` configuration file (".tmux.conf") inside the user's home directory ("/home/ttm4200"):

```
unbind C-b
set-option -g prefix C-a
bind-key C-a send-prefix
```

And to change the command for vertical split from `%` to `\`, horizontal split from `"` to `-`, add the

following lines:

```
bind-key \ split-window -h
bind-key - split-window -v
```

And you can do many interesting customizations such as fast pane-switching, pane-resizing, etc. If you are interested see this [blog](https://www.hamvocke.com/blog/a-guide-to-customizing-your-tmux-conf/) (<https://www.hamvocke.com/blog/a-guide-to-customizing-your-tmux-conf/>). For these changes to take effect, you have to exit from `tmux`, i.e. kill all active sessions. In this course, to change the `tmux` configuration for all containers, you only need to change “`.tmux.conf`” inside the folder from which the image is built (“`ttm4200_base`”, in this case). But you have to rebuild the image and containers for these configuration to take effect.

3.2 Traffic Capture

To capture network traffic, we will use `tcpdump` (<https://www.tcpdump.org/manpages/tcpdump.1.html>), which is a command-line packet capturing utility that allows us to sniff, capture and monitor any type of traffic on a network easily.

- In “container4”, attach to the previously created session:

```
tmux attach -t [session_name]
```

- In one pane, start packet capturing with `tcpdump` and dump the captured packets to a file:

```
sudo tcpdump -i eth0 -w work_dir/container4.pcap
```

- Switch to the other pane `Ctrl+b arrow key`, ping and traceroute “container3” using IPv4 and IPv6 addresses.

```
ping 10.20.30.12
ping -6 fd00::12
traceroute 10.20.30.12
```

We highly recommend getting comfortable with using `tmux`, as it will be worth the time you spent on learning it. Otherwise, you can just open multiple SSH connections

- Stop the `tcpdump` capturing and copy the capture file to your host machine. You can use `scp` (<https://haydenjames.io/linux-securely-copy-files-using-scp/>):

```
scp -P 4000 ttm4200@127.0.0.1:~/work_dir/container4.pcap ./
```

Since the “`ttm4200_base`” directory is mounted on “`container4`” in “`work_dir`”, any file

saved to the “work_dir” inside the container, will be directly saved to the “ttm4200_base” folder

- Using the docker-compose file, stop and remove the containers, as well as the network. You can do so with only one command:

```
docker-compose down
```

Q2. REPORT(10%): Submit your capture file “container4.pcap” along with the report. We will check it and make sure that it is unique. So please don’t copy it from other teams.

3.3 Traffic Analysis

To analyze network traffic, we will use [Wireshark](https://www.wireshark.org/) (<https://www.wireshark.org/>), which is a free and open-source packet analyzer. You can refer to the course’s book in “Chapter 1: Wireshark Lab: Getting Started”, to become familiar with Wireshark.

- In your VM, open the file “container4.pcap” in Wireshark.
- Examine one ping request and its reply using “icmp” filter. Then answer these questions and validate your answer with an **annotated** screenshot from Wireshark.

Q3. REPORT(5%): How long did it take from when the ping request was sent until the reply was received?

Q4. REPORT(5%): Does the ICMP packet have source and destination port numbers?

Q5. REPORT(10%): What are the ICMP type and code numbers in the ping request? What other fields does this ICMP packet have? How many bytes are the checksum, sequence number and identifier fields?

Q6. REPORT(10%): What are the ICMP type and code numbers in the corresponding reply? What other fields does this ICMP packet have? How many bytes are the checksum, sequence number and identifier fields?

Q7. REPORT(10%): Explain why there are ICMP messages: Destination Unreachable (Port unreachable), coming from your destination.

These messages correspond to `traceroute` command. To find the reason, search for the working principles of `traceroute`, especially the value of the used port number

- Examine one ping request and its reply using “icmpv6” filter. Then answer these questions and validate your answer with an **annotated** screenshot from Wireshark.

Q8. REPORT(5%): How long did it take from when the ping request was sent until the reply was received?

Q9. **REPORT(5%):** Does the ICMP packet have source and destination port numbers?

Q10. **REPORT(10%):** What are the ICMP type and code numbers in the ping request? What other fields does this ICMP packet have? How many bytes are the checksum, sequence number and identifier fields?

Q11. **REPORT(10%):** What are the ICMP type and code numbers in the corresponding reply? What other fields does this ICMP packet have? How many bytes are the checksum, sequence number and identifier fields?

4 Optional Exercise

This exercise is to test your understanding of docker. It is not connected to the previous milestones.

- Using docker-compose, build images and run containers from Dockerfile in the directories named “node-server” and “client”. You can complete the following “docker-compose.yml” file to do so:

```
version: '2.3'
networks:
  ttm4200_net:
    enable_ipv6: true
    ipam:
      config:
        #configure subnet range (both IPv6 and IPv4)
        ===== FILL IN HERE =====
services:
  node:
    build: ./node-server/
    container_name: node-server
    hostname: node-server
    cap_add:
      - NET_ADMIN
    ports:
      #map port 8080 on the host machine to port 8080 in the container
      ===== FILL IN HERE =====
    networks:
      ttm4200_net:
        #assign the node service an IPv6 address and IPv4 address
        ===== FILL IN HERE =====
  client:
    build: ./client
    container_name: client
    hostname: client
    user: root
    tty: true
    stdin_open: true
    cap_add:
      - NET_ADMIN
    volumes:
      - ./client/:/home
```



```
networks:
  ttm4200_net:
    #assign the client service an IPv6 address and IPv4 address
    ===== FILL IN HERE =====
```

- Verify that both containers are running, and that a new network has been created.

You will have a network conflict if you have not removed the network in Milestone 3 (i.e. `docker-compose down`), since you are trying to build two networks with the same name.

- Attach to client-container.
- Ping the node-server, using the name you assigned it when you ran the container.
- Ping the node-server again, this time using its IP-address.
- Test the node-server by typing “localhost:8080” in a browser of your choice (**needs to be a browser on the machine which is running Docker**).
- Download the .html-file from the node-server using the “wget” command

```
wget <node-server ip-address>:8080
```

- Edit the “index.html” file located in “/node-server/views” to say “Hello from group <your group ID>”, instead of “Hello World!”.
- Copy the newly edited “index.html” file to the node-server container, and use the path “/app/views/index.html”.

You can either use `docker cp` command or the mounted volume.

- Access “localhost:8080” in the browser.